

Stanford CS336 第 15 讲：中期训练与后训练

标准中文图文课程笔记

五道口纳什 & Codex

2026 年 8 月 30 日



频道：Stanford Online **发布日期：**2026 年 5 月 27 日

时长：1:19:54 **链接：**<https://youtu.be/2oH6PWPrYFo>

字幕：YouTube 平台人工英文字幕 en-US；未使用 Whisper。

目录

1 动机：预训练能力为何还不等于可控助手	3
1.1 本章小结	3
2 证据边界：现代后训练配方并不透明	4
2.1 本章小结	4
3 SFT 数据的演进与共同结构	5
3.1 本章小结	5
4 FLAN：任务模板与混合的力量	6
4.1 本章小结	6
5 Alpaca：蒸馏式指令数据的简洁配方	7
5.1 本章小结	7
6 SFT 目标：只在目标响应上计算交叉熵	8
6.1 本章小结	8
7 数据风格会进入模型，而基准未必能测出	9
7.1 本章小结	9
8 知识抽取、知识注入与幻觉	10
8.1 本章小结	10
9 安全 SFT：从真实交互提取场景	11
9.1 本章小结	11
10 少量高针对性安全数据也能显著改变行为	12
10.1 本章小结	12
11 Tulu 3：把多来源 SFT 组织为可审计管线	13
11.1 本章小结	13
12 把 SFT 数据配方放在一起	14
12.1 本章小结	14
13 中期训练与两阶段训练	15
13.1 本章小结	15
14 RLHF 的三步标准流程	16
14.1 本章小结	16

15 从模仿到优化：生成—验证鸿沟	17
15.1 本章小结	17
16 RLHF 数据：标注者、规范与报酬都是模型的一部分	18
16.1 本章小结	18
17 人口统计与风格偏好：平均标签不等于普遍价值	19
17.1 本章小结	19
18 模型生成反馈与自训练	20
18.1 本章小结	20
19 KL 正则化的偏好优化目标	21
19.1 本章小结	21
20 PPO 在语言模型中的机制	22
20.1 本章小结	22
21 为什么尝试摆脱 PPO	23
21.1 本章小结	23
22 DPO：从 KL 正则化目标消去显式奖励模型	24
22.1 本章小结	24
23 DPO 损失、梯度直觉与代码	25
23.1 本章小结	25
24 DPO 的迭代配方与 PPO 对比	26
24.1 本章小结	26
25 失败模式：过优化、模式坍塌与长度偏差	27
25.1 本章小结	27
26 可复现的中期/后训练协议	28
26.1 本章小结	28
27 总结与延伸	29
27.1 延伸问题	29
27.2 建议练习	29

1 动机：预训练能力为何还不等于可控助手


自回归预训练优化“给定前缀，下一 token 是什么”，却没有显式告诉模型要遵守指令、用何种风格回答、何时拒绝。后训练把潜在能力整理成可由自然语言调用的行为接口。

Instruction following is a remarkable form of control

Prompt: Can you generate a pyplot for the following data: $x = [1, 3, 5, 6, 8]$, $y1 = [2, 3, 5, 18, 1]$, $y2 = [3, 5, 6, 8, 1]$, $y3 = [5, 1, 2, 3, 4]$, $y4 = [9, 7, 2, 3, 1]$. I want $y1, y2$ to be in the same plot, but $y3$ is in another plot next to that plot, $y4$ is in below. I want the legend of $y1$ to be "bob", $y2$ to be "alice", $y3$ to be "bilbo", $y4$ to be "allie". I want the x-axis to be labeled with "time" and y axis to be labeled with "money". I want to add a 10 %-40% random error bar to each curve, through all times (including non-integers). I want smoothed curves to show the plot, and smoothed error bar. Do not use linear interpolation, use smooth interpolation! I want to also add some small zig-zag to the smoothed curve to make it look more real. I want to put a baseline as the mean of each line. I want to put a pie chart below indicating the fraction of the four people in each time step. I also want an animation to show how the fractions are changing in the pie chart in **continuous time**. Interpolate the missing fractions! I also want vertical line animation in other three plots to match the pie chart. I want the fanciest plot. Please add as many fancy things as possible.

Here is a possible code to generate a pyplot for the following data, with some comments to explain the steps:

GPT-4: [We give two snapshots from the animation resulting from the produced code]



Text example from [Bubeck et al 2023]

图 1: 复杂绘图指令被分解并执行，说明指令跟随本身是一种强控制机制。

本讲主线

先用监督微调（SFT）模仿理想响应，再从偏好数据学习“哪个回答更好”；中期训练则在预训练末段改变数据配方，为后训练准备能力与分布。

1.1 本章小结

后训练不是给模型凭空注入全部知识，而是选择、组合并校准预训练已形成的行为。

2 证据边界：现代后训练配方并不透明

公开论文常只给出框架；真实数据来源、标注规范、过滤器与配比通常缺失。旧论文反而可能公开更完整的标注说明。因此复现时必须区分“论文明确披露”“开源实现推断”和“行业传闻”。

Caveat: post-training information is pretty sparse!

Rich information from the pre-chatgpt competition

- Stiennon 2020 (RLHF paper) – contains annotation guidelines, etc
- Bai 2022 (Anthropic HH) – details on how to set up safety annotation, etc

Poor information in modern times..

- Open source: lots of distillation, not too public in open-source release notes
- Closed source: ‘secret sauce’ for data.

In July 2023, a manager ordered workers to study GPT-4's responses closely and figure out why they outperformed Bard's. "Try to come up with feedback that we can share so that experts can write responses better than GPT4 or at least the same," the manager wrote.

Scale AI also created a spreadsheet that compared 1,729 Bard rewrites directly to ChatGPT in October 2023. Each rewrite was rated with labels like "worse than GPT4" or "Needs Some Fixes." In one example, a worker rewrote a Bard review of a nursery chair that managers stamped "worse than GPT4" because it "lacks detail compared to GPT4."

图 2: 课程提醒：现代模型的后训练细节与数据经常不公开。

不要把模型卡当完整配方

只知道数据集名称无法复现结果；还需 chat template、损失掩码、采样权重、去重、过滤、优化器与参考模型版本。

2.1 本章小结

严谨的课程笔记应记录证据等级，避免把猜测写成确定事实。

3 SFT 数据的演进与共同结构

开放 SFT 从 FLAN、Self-Instruct 发展到 Alpaca、ShareGPT/Vicuna、OpenAssistant、WizardLM、Tulu 与 Nemotron。名称不同，训练记录通常都能规约为“上下文/指令 x 与目标响应 y ”。




图 3: 开放世界中 SFT 数据从任务模板、自动生成到真实对话和复杂后训练配方的演进。

三类来源

人工编写的任务与答案质量高但昂贵；真实对话自然但许可与隐私复杂；教师模型生成可扩展，但会继承教师风格与错误。

3.1 本章小结

数据集的关键差异不只在规模，更在任务覆盖、响应来源、风格、许可和验证方式。

4 FLAN：任务模板与混合的力量

FLAN 把大量 NLP 数据集转成自然语言指令，并用多个模板表述同一任务。它证明了任务混合与措辞多样性可以显著增强零样本迁移，但随机样例也暴露出输入长短和输出格式差异巨大。

FLAN – random examples	
Stephanie - Can you finalize the attached and have it signed. I need to initial it, but it needs to be signed by Brad Richter. Thanks. Write a subject line for this email.	Ronald Chisholm LOI
Ahold to Sell Spain Operations to Permira (AP) AP - The Dutch supermarket retailer Ahold, seeking to streamline global operations and reduce debt, said Sunday it will sell its holdings in Spain to Permira Funds for about #36;849 million. What is this text about? OPTIONS: - World - Sports - Business - Science/Tech	Business
Write highlights for this article: Sauntering down leafy avenues past typical Dutch step-gabled buildings, carpets of crocuses, the occasional cyclist whooshing quietly by and the sun bouncing off the canal, I can think of few more pleasant places to spend a spring weekend than The Hague. [...]The most prized Dutch colony was apparently run from The Hague (aka 'the Widow of Indonesia'). For summer: the excellent Grapes & Olives and the Café de Bieb on Veenkade will have boats on a newly renovated bit of canal, with music and boat trips planned, including a tour at the Kompaan brewery (I can vouch for the product). My art gorging finishes at the Panorama, an extraordinary, 14 ft-high circular panorama of The Hague's seaside resort, Scheveningen, painted by Hendrik Mesdag. I take a 15-minute tram to the place itself and gorge on scallops and sole at Catch by Simonis on the harbour front. The Hague has netted a new fan. British Airways (0844 493 0787, www.ba.com) flies to Rotterdam Den Haag (half-hour taxi) from £100 return. For tours with Remco Dörr, email remco.dorr@hotmail.com. Double rooms at the Carlton Ambassador (00 31 70 363 0363, www.carlton.nl/ambassador) cost from £84, room only. More information on travel in The Netherlands at www.holland.com.	The least known of the Dutch cities, The Hague was a village until 1806. It owes its growth to Louis Bonaparte, Napoleon's brother, who ruled here. The city has a wealth of art, including Vermeer's 'Girl With A Pearl Earring'
Here is some data about a restaurant: name = Aromi, eatType = coffee shop, food = English, customer rating = 5 out of 5, area = city centre. Write a sentence that includes the following data about a restaurant	In the city centre there is a coffee shop with a customer rating of 5 out of 5 called Aromi which serves English food.

图 4: FLAN 随机样例：分类、摘要与结构化信息抽取可共享统一的文本接口。

模板泄漏

同一基准的训练模板与测试模板过近，会把“遵循新指令”混同于“识别熟悉模板”；划分应按原任务族和来源去重。

4.1 本章小结

模板把异构任务统一起来，但模板多样性与跨划分去重决定泛化是否真实。

5 Alpaca：蒸馏式指令数据的简洁配方

Alpaca 用强模型根据种子任务扩写指令—响应对，使小团队也能构造大规模 SFT 数据。其优势是便宜、格式统一，代价是响应同质化、事实错误和教师偏差。

Alpaca – random examples	
Give three tips for staying healthy.	1.Eat a balanced diet and make sure to include plenty of fruits and vegetables. 2. Exercise regularly to keep your body active and strong. 3. Get enough sleep and maintain a consistent sleep schedule.
What does “algorithm” mean?	An algorithm is a step-by-step set of instructions used to solve a problem or achieve a goal. It is a process of breaking down complex tasks into smaller, more manageable steps. Algorithms can be used to find the shortest path from one place to another, sort data, or render data from an immense database in a few seconds.
Find the average number in a list.	def avg_list(nums): total = 0 for n in nums: total += n return total / len(nums) print(avg_list([4, 3, 6, 10, 8])) # Output 6.4

图 5: Alpaca 随机样例：健康建议、概念解释与代码任务被整理为统一指令—回答对。

Listing 1: SFT 记录的最小规范化结构

```
1 record = {
2     "messages": [
3         {"role": "user", "content": instruction},
4         {"role": "assistant", "content": response},
5     ],
6     "source": "dataset/version",
7     "teacher": teacher_id,
8     "license": license_id,
9 }
```

5.1 本章小结

合成指令数据应携带教师、提示词、采样参数和验证结果，不能只留下最终文本。

6 SFT 目标：只在目标响应上计算交叉熵

令输入/上下文为 x ，目标响应 token 为 $y_1, \dots, y_T$ ，监督微调损失为

$$\mathcal{L}_{\text{SFT}}(\theta) = - \sum_{t=1}^T m_t \log \pi_{\theta}(y_t \mid x, y_{<t}).$$

θ 是模型参数； π_{θ} 是条件 token 分布； $y_{<t}$ 是此前目标 token； $m_t \in \{0, 1\}$ 是损失掩码。通常用户和 system token 的 $m_t = 0$ ，assistant 目标 token 的 $m_t = 1$ 。

Listing 2: 带响应掩码的 SFT 损失

```
1 logits = model(input_ids).logits[:, :-1]
2 targets = input_ids[:, 1:]
3 loss = cross_entropy(logits.transpose(1, 2), targets,
4                       reduction="none")
5 loss = (loss * assistant_mask[:, 1:]).sum()
6 loss /= assistant_mask[:, 1:].sum().clamp_min(1)
```

常见错误

错误的 chat template 或边界 token 会让掩码错位；训练前应反解若干样本，逐 token 检查角色、特殊符号和 mask。

6.1 本章小结

SFT 在数学上仍是最大似然，但“对哪些 token 求似然”决定模型学习谁的行为。

7 数据风格会进入模型，而基准未必能测出

不同 SFT 数据在礼貌程度、拒答、列表偏好、长度和语气上差异明显。传统准确率基准可能几乎不变，但用户偏好会显著变化。

What about benchmarks?

These factors are (mostly) not that relevant for other benchmark perfs

Table 3: Comparison of different instruction tuning datasets, showing that different instruction-tuning datasets can excel in different aspects, and mixtures perform best on average. Cells are blue if the finetuning boosts the vanilla LLaMA performance, and orange if the finetuning hurts the performance.

	MMLU (factuality)	GSM (reasoning)	BBH (reasoning)	TyDiQA (multilinguality)	Codex-Eval (coding)	AlpacaEval (open-ended)	Average
	EM (0-shot)	EM (8-shot, CoT)	EM (3-shot, CoT)	F1 (1-shot, GP)	P@10 (0-shot)	Win % vs Davinci-003	
Vanilla LLaMa 13B	42.3	14.5	39.3	43.2	28.6	-	-
+SuperNI	49.7	4.0	4.5	50.2	12.9	4.2	20.9
+CoT	44.2	40.0	41.9	47.8	23.7	6.0	33.9
+Flan V2	50.6	20.0	40.8	47.2	16.8	3.2	29.8
+Dolly	45.6	18.0	28.4	46.5	31.0	13.7	30.5
+Open Assistant 1	43.3	15.0	39.6	33.4	31.9	58.1	36.9
+Self-instruct	30.4	11.0	30.7	41.3	12.5	5.0	21.8
+Unnatural Instructions	46.4	8.0	33.7	40.9	23.9	8.4	26.9
+Alpaca	45.0	9.5	36.6	31.1	29.9	21.9	29.0
+Code-Alpaca	42.5	13.5	35.6	38.9	34.2	15.8	30.1
+GPT4-Alpaca	46.9	16.5	38.8	23.5	36.6	63.1	37.6
+Baize	43.7	10.0	38.7	33.6	28.7	21.9	29.4
+ShareGPT	49.3	27.0	40.4	30.5	34.1	70.5	42.0
+Human data mix.	50.2	38.5	39.6	47.0	25.0	35.0	39.2
+Human+GPT data mix.	49.3	40.5	43.3	45.6	35.9	56.5	45.2

图 6: 多种指令微调数据在不同基准上的表现不一致，平均分掩盖了任务间差异。

双轨评估

同时报告能力基准与行为评估：事实/推理是否正确，以及风格、帮助性、安全性和长度是否符合目标。

7.1 本章小结

单一平均分不能回答“数据是否更好”；评价维度必须覆盖产品真正关心的行为。

8 知识抽取、知识注入与幻觉

SFT 擅长让模型调用已经学会的知识，却不一定适合可靠注入新事实。若目标事实超出模型已有能力，行为克隆可能先提高训练集拟合，随后损害泛化并增加幻觉。

Knowledge extraction and alignment


Folklore: Fine-tuning a model on ‘facts it doesn’t know’ makes it hallucinate

Hallucination and Behavior Cloning

- Conceptual model of fine-tuning neural nets for question answering
- Neural net has “knowledge graph” stored in weights, with confidence level on each edge
- Small scale fine-tuning learns a simple function that operates on the knowledge graph and outputs token predictions
 - e.g. Q: what is the genre of Star Wars. A: Sci-Fi
- If you clone on correct answers that aren’t in the knowledge graph, you’re teaching the net to hallucinate
- Suppose labeler knows about spinoff film, but net’s knowledge graph doesn’t ...
- Clone on: Q: what was the name of spinoff film centering on Han Solo. A: Solo
- If you clone on incorrect answers that are in knowledge graph, then you’re training net to withhold information
- Suppose labeler doesn’t know about Tom Simpson but net does, trains with “I don’t know”
- Behavior cloning target should depend on network’s knowledge (which is unknown to experimenter)
- Models that are trained using targets computed by another agent will always have hallucination problem



[Schulman 2023]



[Gekhman 2023]

图 7: 知识抽取与对齐：训练已知与未知事实会呈现不同学习轨迹。

实践判断

若目标是“教会回答格式与检索策略”，SFT 很合适；若目标是大量新知识，应优先考虑预训练/中期训练、检索增强，或可更新的外部工具。

8.1 本章小结

把行为对齐与事实记忆分开设计，能减少把流畅模仿误当成可靠知识获得。

9 安全 SFT：从真实交互提取场景

安全数据的难点不是写几句拒答，而是覆盖真实用户如何表达恶意、误导、诈骗、危险操作和边界请求。WildChat 等真实对话可帮助发现长尾场景，再由规则与人工审查转成训练任务。




图 8: 从真实用户对话中抽取安全场景、类别和对抗性提示。

隐私与二次伤害

真实对话必须做个人信息去除、许可检查和危险内容访问控制；审查流程也要保护标注者。

9.1 本章小结


安全覆盖来自系统化场景发现，不来自单一“拒绝模板”。

10 少量高针对性安全数据也能显著改变行为

课程给出的实验显示，约数百条 Alpaca 风格安全样例即可明显改善多类安全指标。这说明模型往往已经具备相关能力，只需少量数据指定决策边界和表达方式。

Safety-tuning with just a little data

Significant improvements to safety with ~500 samples.



Adding 500 Alpaca-style examples makes models follow safety guidelines

🔍 🔍 🔍 🔍 🔍 🔍 🔍 🔍 🔍 🔍

图 9: 约 500 条安全样例即可在多个数据集上显著降低不安全得分。

不是越多越好

安全样例比例过高可能造成过度拒答；应联合测量攻击成功率、正常请求帮助率与拒答校准。

10.1 本章小结

高杠杆数据的价值来自边界覆盖与标签一致，而不只是样本数量。

11 Tulu 3：把多来源 SFT 组织为可审计管线

Tulu 3 展示了开放后训练中较完整的配方：通用指令、数学/代码、工具、对话与安全数据分别构建，再经筛选、格式化与配比进入训练。




图 10: Tulu 3 的公开数据表给出安全与非服从数据的来源及规模。

Listing 3: 多来源 SFT 混合与审计

```
1 for step in range(num_steps):
2     source = rng.choice(sources, p=mixture_weights)
3     batch = loaders[source].next()
4     loss = sft_step(batch)
5     audit.log(step=step, source=source,
6               tokens=batch.assistant_tokens, loss=loss)
```

11.1 本章小结

开放配方的贡献是让数据来源、规模和阶段可追踪，而不是只发布一个合并后的 JSON。

12 把 SFT 数据配方放在一起

课程归纳出三点：SFT 最擅长抽取预训练已有行为；即使事实正确的数据也可能伤害模型；安全、指令跟随和风格等正确行为的小样本能带来大变化，但总体仍受长期数据收益影响。

Putting it together – SFT Data

1. Instruction fine-tuning (SFT) works best when we are just extracting pre-training behaviors, not adding new ones
2. Adding (factually correct!) data can sometimes hurt
3. Small amounts of the right kinds of behavior (safety, instruction-following, style) make a big difference, but there is a long-tail that benefits from more data



图 11: SFT 数据部分的三条结论：能力抽取、事实数据风险与小量行为数据的高杠杆效应。

配方单位

建议按“能力簇”管理数据：基础对话、知识问答、推理、代码、工具、安全与风格；每簇记录独立规模、采样率、epoch 和评估集。

12.1 本章小结


SFT 配方是一组目标行为的预算分配，不是把所有公开指令集简单拼接。

13 中期训练与两阶段训练

“midtraining”或“两阶段训练”在预训练后段更换数据分布：稳定阶段使用宽泛混合，衰减阶段提高高质量、代码、数学或指令型数据占比，并配合学习率衰减。

‘Midtraining’ / ‘Two-phase training’

The recipe is common knowledge among many LLM companies (but not documented)



Widely used by most models today..

Publicized in recent Chinese-derived LMs (miniCPM, jetMoE)

图 12: 稳定阶段与衰减阶段的数据混合饼图展示了两阶段训练的配方变化。

若第 k 阶段数据分布为 $p_k(d)$ 、学习率为 η_k ，更新可写为

$$\theta_{t+1} = \theta_t - \eta_k \mathbb{E}_{d \sim p_k} [\nabla_{\theta} \ell(d; \theta_t)].$$

d 是训练样本， p_k 是阶段混合， ℓ 是 token 损失。数据分布与学习率同时改变，消融时必须分别控制。

13.1 本章小结

中期训练连接“广泛世界建模”与“目标能力塑形”，它仍是 token 级预训练而非偏好优化。

14 RLHF 的三步标准流程

典型流程是：先收集示范做 SFT；再对同一提示的多个回答做成对排序并训练奖励模型；最后用强化学习优化策略，同时约束其不要偏离参考策略。




图 13: InstructGPT 风格流程：示范、比较、奖励建模与策略优化。

数据与算法不可分

偏好数据决定奖励信号，优化算法决定模型能多激进地追逐它；高质量比较标签也不能抵消失控的过优化。

14.1 本章小结

RLHF 的核心资产不只是算法，还包括提示分布、候选生成策略、标注规范与人群。

15 从模仿到优化：生成—验证鸿沟

SFT 拟合参考答案分布 $p^*(y|x)$ ；RLHF 则寻找使可测奖励最大的策略。人们未必能直接写出最佳答案，却常能在候选中判断哪个更好，这就是生成—验证鸿沟。

From imitation to optimization

Imitation (SFT)

Fit $\hat{p}(y|x) \approx p^*(y|x)$ for some reference distribution $p^*(y|x)$

- Pure generative modeling perspective
- Requires samples from reference policy

Optimization (RLHF)

Find $\hat{p}(y|x)$ such that $\max_p E_p[R(y, x)]$ for a reward $R(y, x)$



- Maximize some reward function that we can measure
- LMs are policies, not a model of some distribution

🔍 🔍 🔍 🔍 🔍 🔍 🔍 🔍 🔍 🔍

图 14: 模仿拟合参考分布，优化则直接提高可测奖励。

Why optimize? G-V gap

People don't always write the thing that they prefer in LM outputs..



From: Zhang et al – "Benchmarking Large Language Models for News Summarization" 2023

图 15: 不同标注者对摘要风格偏好差异明显，说明验证也并非绝对客观。

15.1 本章小结

偏好比较通常比从零写答案容易，但“容易比较”不等于“没有主观性”。

16 RLHF 数据：标注者、规范与报酬都是模型的一部分

偏好标签由具体人群在具体指南与工资条件下产生。专家标注成本高，但复杂领域的通才标签可能不可靠；低报酬又会诱导匆忙判断。


Large variation in compensation

Lots of growth in “expert annotation”

Under Project Stagecraft, as it is internally known at [data-labeling startup Handshake AI](#), freelancers are being paid at least \$50 an hour to create materials that will be used by OpenAI’s ChatGPT to understand various occupations, one contractor told Business Insider.

Business Insider reviewed dozens of pages of documents for the project, including step-by-step instructions for contractors. In these files, contractors from various industries were directed to develop a persona and create tasks that reflect real work.

The project, which has employed about 3,000 to 4,000 freelancers, is ongoing, the contractor said.



<https://www.businessinsider.com/ai-data-labeling-annotators-pay-subject-experts-generalists-gig-workers-2025-12>

图 16: 专家标注与通用标注报酬差异巨大，直接影响可扩展性和标签质量。

最小标注元数据

匿名标注者组别、资格、指南版本、任务耗时、候选顺序随机化、平局/跳过、复核结果与报酬区间。

16.1 本章小结

标注平台不是中性传感器；人员选择与激励塑造了奖励函数。

17 人口统计与风格偏好：平均标签不等于普遍价值

不同宗教、文化、年龄和教育背景可能对同一回答给出不同判断。把标签平均成一个奖励分数，会把群体差异压缩成单一“标准用户”。

RLHF and data - demographics

The annotator distribution for RLHF can significantly shift its behaviors

What gender do you identify as?		AI21				OpenAI			
		j1-grande	j1-jumbo	j1-grande-v2-beta	ada	gpt-4o	gpt-4o-mini	gpt-4o-mini-2024-07-18	gpt-4o-mini-2024-08-07
Male		50.0%	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%	50.0%
Female		44.4%	44.4%	44.4%	44.4%	44.4%	44.4%	44.4%	44.4%
Nonbinary / other		5.6%	5.6%	5.6%	5.6%	5.6%	5.6%	5.6%	5.6%
What ethnicity do you identify as?									
White / Caucasian		31.6%	31.6%	31.6%	31.6%	31.6%	31.6%	31.6%	31.6%
Southeast Asian		52.6%	52.6%	52.6%	52.6%	52.6%	52.6%	52.6%	52.6%
Indigenous / Native American / Alaskan Native		0.0%	0.0%	0.0%	0.0%	0.0%	0.0%	0.0%	0.0%
East Asian		5.3%	5.3%	5.3%	5.3%	5.3%	5.3%	5.3%	5.3%
Middle Eastern		0.0%	0.0%	0.0%	0.0%	0.0%	0.0%	0.0%	0.0%
Latino		15.8%	15.8%	15.8%	15.8%	15.8%	15.8%	15.8%	15.8%
Black / of African descent		10.5%	10.5%	10.5%	10.5%	10.5%	10.5%	10.5%	10.5%
What is your nationality?									
Filipino		22%	22%	22%	22%	22%	22%	22%	22%
Bangladeshi		22%	22%	22%	22%	22%	22%	22%	22%
American		17%	17%	17%	17%	17%	17%	17%	17%
Albanian		5%	5%	5%	5%	5%	5%	5%	5%
Brazilian		5%	5%	5%	5%	5%	5%	5%	5%
Canadian		5%	5%	5%	5%	5%	5%	5%	5%
Colombian		5%	5%	5%	5%	5%	5%	5%	5%
Indian		5%	5%	5%	5%	5%	5%	5%	5%
Uruguayan		5%	5%	5%	5%	5%	5%	5%	5%
Zimbabwean		5%	5%	5%	5%	5%	5%	5%	5%
What is your age?									
18-24		26.3%	26.3%	26.3%	26.3%	26.3%	26.3%	26.3%	26.3%
25-34		47.4%	47.4%	47.4%	47.4%	47.4%	47.4%	47.4%	47.4%
35-44		10.5%	10.5%	10.5%	10.5%	10.5%	10.5%	10.5%	10.5%
45-54		10.5%	10.5%	10.5%	10.5%	10.5%	10.5%	10.5%	10.5%
55-64		5.3%	5.3%	5.3%	5.3%	5.3%	5.3%	5.3%	5.3%
65+		0%	0%	0%	0%	0%	0%	0%	0%
What is your highest attained level of education?									
Less than high school degree		0%	0%	0%	0%	0%	0%	0%	0%
High school degree		10.5%	10.5%	10.5%	10.5%	10.5%	10.5%	10.5%	10.5%
Undergraduate degree		52.6%	52.6%	52.6%	52.6%	52.6%	52.6%	52.6%	52.6%
Master's degree		36.8%	36.8%	36.8%	36.8%	36.8%	36.8%	36.8%	36.8%
Doctorate degree		0%	0%	0%	0%	0%	0%	0%	0%

[Santurkar+ 2023]

图 17: 标注者人口统计分布变化会显著改变不同模型行为的评分。

多数偏好并非安全真理

应分组报告一致率，保留分歧与平局；对安全和公平等高风险决策，需由政策约束而非简单多数投票决定。

17.1 本章小结


RLHF 对齐的是被采样、被表达和被聚合的偏好，而不是抽象的“人类价值”。

18 模型生成反馈与自训练

AI feedback 可生成候选、批注偏好或充当裁判，显著降低成本；UltraFeedback、Zephyr、Tulu 等配方体现了这一趋势。但同一家族模型生成、筛选并评估会形成闭环偏差。

RLHF and data - LM-generated

For people following the frontier – AI feedback often used for RLHF



Used in e.g. Olmo, Zephyr, etc

<https://thesequence.substack.com/p/the-sequence-chat-hugging-faces-lewis>

图 18: 模型生成偏好反馈：候选生成、偏好标注与训练形成自动化管线。

Listing 4: 带独立验证的模型反馈循环

```
1 candidates = policy.sample(prompt, n=K)
2 ranking = judge.rank(prompt, candidates)
3 if verifier.agrees(ranking) and not leakage(prompt):
4     preference_store.add(prompt, ranking, judge.version)
5 else:
6     human_queue.add(prompt, candidates)
```

18.1 本章小结

自训练的瓶颈从“生成数据”转为“获得独立、校准且不循环自证的验证信号”。

19 KL 正则化的偏好优化目标

给定提示分布 $x \sim \mathcal{D}$ ，策略 π_θ 的经典目标是

$$\max_{\theta} \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta(\cdot|x)}[r_\phi(x, y)] - \beta \mathbb{E}_{x \sim \mathcal{D}}[\text{KL}(\pi_\theta(\cdot|x) \parallel \pi_{\text{ref}}(\cdot|x))].$$

r_ϕ 是奖励模型； π_{ref} 是通常由 SFT 得到的参考策略； $\beta > 0$ 控制保守程度；KL 项惩罚策略偏离参考分布。

解释

第一项鼓励高奖励，第二项防止模型利用奖励模型漏洞并保持语言质量。 β 太大几乎不学习，太小则易奖励黑客和模式坍塌。

19.1 本章小结

RLHF 是带信赖区域的奖励优化，而不是无限最大化一个不完美评分器。

20 PPO 在语言模型中的机制

PPO 使用旧策略采样 on-policy 序列，并限制新旧策略概率比变化。令

$$\rho_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)},$$

其中 $s_t = (x, y_{<t})$, $a_t = y_{t\circ}$ 裁剪代理目标为

$$L^{\text{clip}} = \mathbb{E}_t [\min(\rho_t A_t, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) A_t)] .$$

A_t 是优势估计， ϵ 是裁剪半径；min 防止通过过大概率更新获取虚假收益。

PPO in language modeling

From InstructGPT

Reinforcement learning (RL). Once again following Stiennon et al. (2020), we fine-tuned the SFT model on our environment using PPO (Schulman et al., 2017). The environment is a bandit environment which presents a random customer prompt and expects a response to the prompt. Given the prompt and response, it produces a reward determined by the reward model and ends the episode. In addition, we add a per-token KL penalty from the SFT model at each token to mitigate over-optimization of the reward model. The value function is initialized from the RM. We call these models “PPO.”

We also experiment with mixing the pretraining gradients into the PPO gradients, in order to fix the performance regressions on public NLP datasets. We call these models “PPO-ptx.” We maximize the following combined objective function in RL training:

$$\begin{aligned} \text{objective}(\phi) = & E_{(x,y) \sim D_{\text{SFT}}} [r_\theta(x,y) - \beta \log(\pi_\phi^{\text{RL}}(y|x)/\pi^{\text{SFT}}(y|x))] + \\ & \gamma E_{x \sim D_{\text{pretrain}}} [\log(\pi_\phi^{\text{RL}}(x))] \end{aligned} \quad (2)$$

where π_ϕ^{RL} is the learned RL policy, π^{SFT} is the supervised trained model, and D_{pretrain} is the pretraining distribution. The KL reward coefficient, β , and the pretraining loss coefficient, γ , control the strength of the KL penalty and pretraining gradients respectively. For “PPO” models, γ is set to 0. Unless otherwise specified, in this paper InstructGPT refers to the PPO-ptx models.

..this is very innocuous looking

图 19: InstructGPT 中 PPO 目标同时包含奖励、KL 约束与预训练梯度混合。

Listing 5: PPO-RLHF 的概念级训练循环

```
1 responses = old_policy.generate(prompts)
2 rewards = reward_model(prompts, responses)
3 advantages = estimate_advantage(rewards, value_model)
4 ratio = exp(policy.logp(responses) - old_policy.logp(responses))
5 objective = minimum(ratio * advantages,
6                     clip(ratio, 1-eps, 1+eps) * advantages)
7 loss = -objective.mean() + beta * kl_to_reference(policy)
```

20.1 本章小结

PPO 功能强但系统复杂：策略、旧策略、参考策略、奖励模型与价值模型必须协调。

21 为什么尝试摆脱 PPO

PPO 的 on-policy 采样成本高、实现细节多且训练不稳定。课程回顾了控制 token、只训练胜者、奖励模型筛选 best-of- N 等简化尝试；它们部分有效，却未充分利用输家与参考策略信息。

Can we get rid of PPO?

Question for many: can we avoid doing any 'RL' ? (i.e. on-policy RL algorithms)

Some reasonable stuff people thought about

- Train the model with a control token – (SFT on the pairs, prepend [GOOD] to chosen, [BAD] to not chosen)
- Train the model on only preferred output
- Train a reward model, get LM outputs, train on the preferred output
- Train a reward model, get 1024 LM outputs, take the best one.

图 20: 摆脱 PPO 的早期思路：控制 token、只学偏好回答、奖励筛选与 best-of- N 。

best-of- N 的隐性代价

N 越大越可能找到奖励模型的漏洞，同时推理成本线性增加；必须用独立评估检查真实质量。

21.1 本章小结

简化算法的目标不是少写代码，而是减少在线组件并保留对偏好差异的有效学习信号。

22 DPO: 从 KL 正则化目标消去显式奖励模型

在非参数策略假设下, KL 正则化目标的最优策略满足

$$\pi^*(y|x) = \frac{1}{Z(x)} \pi_{\text{ref}}(y|x) \exp\left(\frac{r(x, y)}{\beta}\right),$$

其中 $Z(x)$ 是对所有回答归一化的配分函数。反解奖励:

$$r(x, y) = \beta \log \frac{\pi^*(y|x)}{\pi_{\text{ref}}(y|x)} + \beta \log Z(x).$$

在同一提示的胜者与败者差分中, $\log Z(x)$ 抵消, 于是无需显式拟合奖励模型。

DPO – RLHF without tears?

Try to simplify PPO by..

- Getting rid of the reward model
- Getting rid of any on-policy stuff (rollouts, outer loops etc)

Instead

- Take gradient steps on log-loss of good stuff
- Take negative gradient steps on bad stuff (appropriately weighted).




图 21: DPO 从 KL 正则化 RLHF 的闭式最优策略反解隐式奖励。

22.1 本章小结

DPO 的关键是用策略相对参考策略的对数概率比表示奖励差。

23 DPO 损失、梯度直觉与代码

对偏好对 (x, y_w, y_l) ， y_w 是胜者， y_l 是败者，DPO 损失为

$$\mathcal{L}_{\text{DPO}} = -\log \sigma \left(\beta \left[\log \frac{\pi_{\theta}(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \log \frac{\pi_{\theta}(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right] \right).$$

σ 是 sigmoid； β 控制相对参考策略的变化尺度。梯度提高胜者相对概率、降低败者相对概率；模型已强烈区分时，更新自动变小。

Listing 6: DPO 批次损失

```
1 win_margin = policy_logp_win - ref_logp_win
2 lose_margin = policy_logp_lose - ref_logp_lose
3 logits = beta * (win_margin - lose_margin)
4 loss = -logsigmoid(logits).mean()
```

与朴素正/负 SFT 的区别

参考策略比值校准了模型原先的偏好；sigmoid 权重又让“已经分对”的样本少更新，从而避免所有偏好对使用相同步长。


23.1 本章小结

DPO 像分类式 SFT，但其目标来自偏好概率模型与 KL 正则化推导。

24 DPO 的迭代配方与 PPO 对比

DPO 可以离线训练，也可循环“采样候选—收集偏好—更新策略”。开放模型常组合奖励筛选、专家偏好数据、SFT 与多轮 DPO。

DPO in LLaMA (and other open models)



DPO + Expert iteration for post-training

图 22: 开放模型中的 DPO 与专家迭代：候选生成、偏好标注、SFT 和 DPO 交替进行。

	PPO	DPO
数据	训练中在线采样	固定或周期刷新偏好对
组件	策略、价值、奖励、参考、旧策略	策略与冻结参考策略
优点	可直接优化序列奖励	简单稳定，像监督训练
风险	不稳定、昂贵、奖励黑客	离线分布受限、偏好数据偏差

24.1 本章小结


DPO 降低系统复杂度，但没有消除偏好采集、分布漂移和过优化问题。

25 失败模式：过优化、模式坍塌与长度偏差

优化不完美奖励时，早期真实质量上升，随后模型开始利用评分器漏洞；同时输出分布可能坍塌，熵下降，表达越来越单一。长度也常成为奖励代理，导致冗长回答获得不成比例的优势。

Things to watch out for in RLHF

Overoptimization / overfitting on the reward



Mode collapse / entropy




图 23: RLHF 风险：奖励过拟合与模式坍塌可在继续优化时恶化。

监控三条曲线

同时跟踪训练奖励、独立人工/强验证器质量和到参考策略的 KL；若奖励继续上升而独立质量下降，应立即早停并审计奖励漏洞。

25.1 本章小结

训练目标的持续改善并不保证真实目标改善，独立验证与保守更新是必要条件。

26 可复现的中期/后训练协议

1. 冻结预训练 checkpoint、tokenizer 与 chat template；
2. 按能力簇记录 SFT 来源、许可、过滤、配比与实际 token；
3. 对用户/system/assistant token 逐项验证损失掩码；
4. 中期训练单独记录阶段边界、混合与学习率；
5. 偏好数据记录候选策略、采样参数、标注者与指南；
6. PPO/DPO 都保留参考策略哈希与 β ；
7. 以固定预算比较 SFT、DPO、PPO，并报告能力、安全、风格、KL 与长度；
8. 失败样例进入新的数据/评估集，但防止训练—测试污染。

Listing 7: 后训练运行清单的机器可读骨架

```
1 run = dict(  
2     base_checkpoint="sha256:...", tokenizer="sha256:...",  
3     chat_template="v3", sft_manifest="sft-v7.jsonl",  
4     preference_manifest="pref-v4.jsonl", algorithm="dpo",  
5     beta=0.1, seed=2026, evaluator="independent-v2")  
6 assert no_overlap(run["sft_manifest"], heldout_eval)
```

26.1 本章小结

可复现性来自 checkpoint、数据、格式、标签与算法状态的联合版本化。

27 总结与延伸

Recap of the lecture

RLHF recap:

1. RLHF data collection is (also) hard! Many confounding factors
2. RLHF algorithms are a bit more complex than SFT – esp. PPO
3. Be mindful of the impact of (over) optimizing for rewards

图 24: 课程回顾: RLHF 数据收集困难, 算法比 SFT 复杂, 并需警惕奖励过优化。

最终结论

SFT 决定模型“像什么样的助手”，偏好优化决定它在候选行为中“更偏向哪一个”，中期训练则提前塑造可被后训练调用的能力。三者都首先是数据分布设计问题，其次才是算法问题。

27.1 延伸问题

研究可继续追问：如何用多群体偏好而非单一平均奖励；如何发现奖励黑客而非事后修补；怎样动态选择 β ；何时 DPO 的离线偏好分布不足，需要在线策略改进；中期数据变化与后训练收益是否具有可预测的 scaling law。

27.2 建议练习

1. 在小模型上逐 token 验证 SFT mask；
2. 用合成偏好实现 DPO，并画胜负 margin；
3. 扫描 β ，比较偏好胜率与 KL；
4. 构造长度偏置奖励，观察过优化；
5. 对同一偏好集按标注者群体重算结果。